

Homework Report #4: Time Series

Author: Isaac Han

Email: cogitoergosum01001@gmail.com

Class: Data Analysis: Statistical Modeling and Computation in applications

Professors: Prof. Uhler, Prof. Jegelka

Table of Contents

DATA:	1
Information about the data set	1
Data Parameters	2
Pre-Processing	2
Partitioning	3
Detrending: Quadratic fitting	3
Written Exercise:	4
Q1: Plot the periodic signal for each month	4
Q2: Plot the final fit	5
Q3: Reported error of RMSE and MAPE(Max 200 words)	6
Q4: Ratio of values	7
Written Exercise: ARIMA	10
Q1: Consider MA(1) model, find autocovariance function	10
Q2: AR(1) model	13
Written Exercise: CPI and BER	15
Data Explanation:	16
Data loading	16
Q0: Verify Your Result By plotting(As recommended by Alan)	16
Q1: Repeat the model fitting and evaluation procedure, with rate	17
1: Description of calculation	17
Q2: Describe how the data has been detrended and plot	18
Q3: Statement of and justification for the chosen AR(p) model	20
Q4: Describe the final model;	21
Q5: Which AR(p) model to choose	23
Q6: Overlay your estimates	24
Written Assignment #3	27
Q1: Plot the cross correlation function between the and inflation rate	27
Q2: Fit SARIMAX	27
Q3: RMSE	28
Q4: Improving Your Model	28

DATA:

Information about the data set

Named as C02.csv, it provides concentration of CO₂ in ppm unit, derived from in situ air measurements, at Mauna Loa, Observatory, Hawaii. The observatory is located at Latitude of 19.5° North and Longitude of 155.6° W with elevation of 3397m.

The concentration of CO₂ was recorded each month starting from March 1958. And only the concentration given in column 5 was used as instructed by the course. The data set contains NA(Not A Number Type) and appropriate data processing was required.

Data Parameters

- $C_i = F(t_i) + P_i + R_i$ where $F : \rightarrow F(t)$ accounts for the long-term trend
- t_i is time at the middle of the i th month, measured in fractions of years after Jan 15, 1958. $t_i = \frac{i + 0.5}{12}$
- P_i is periodic in i with a fixed period, accounting for the seasonal pattern.
- R_i is the remaining residual that accounts for all other influences.
- Data and test sets are split into 80:20


Pre-Processing

```
opts = detectImportOptions('C02.csv');

T = readtable('C02.csv',opts);
T = T(3:end,:);
dim = size(T);
T.count = (0:dim(1)-1)';
T.time = (T.count + 0.5)/12;

processed_T = T(T.C02 ~= -99.99,:);
any(processed_T.C02 == -99.99);
writetable(processed_T, 'C02Processed.csv');
writetable(T, 'C02notprocessed.csv');
clear
```

```
opts = detectImportOptions('C02Processed.csv');
% preview('C02Processed.csv',opts)
T = readtable('C02processed.csv');
plot(T.time, T.C02)
ylim([300 500])
```



Partitioning

```
cutoff = round(size(T,1)*0.8)
```

```
cutoff = 587
```

```
training = {T.time(1:cutoff), T.CO2(1:cutoff)};
test = {T.time(cutoff+1:end), T.CO2(cutoff+1:end)};
```

Detrending: Quadratic fitting

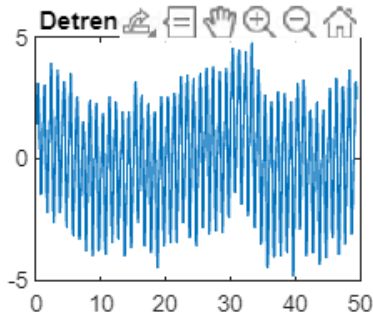
```
x = [ones(length(training{1}),1),training{1},training{1}.^2];
y = training{2}
```

```
y = 587×1
    315.7000
    317.4500
    317.5100
    315.8600
    314.9300
    313.2100
    313.3300
    314.6700
    315.5800
    316.4800
        ⋮
```

```
b = x\y
```

```
b = 3×1
    314.1006
         0.8021
         0.0121
```

```
T_cutoff = T(1:cutoff, :);
T_cutoff.detrend = y - x*b;
plot(T_cutoff.time,T_cutoff.detrend)
title('Detrended(Quadratic) Data')
xlim([0 50])
```



Written Exercise:

Q1: Plot the periodic signal for each month

From previous section, detrending,

$$\begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix} = \begin{bmatrix} 1 & x_1 & x_1^2 \\ 1 & x_2 & x_2^2 \\ \vdots & \vdots & \vdots \\ 1 & x_n & x_n^2 \end{bmatrix} \begin{bmatrix} \beta_0 \\ \beta_1 \\ \beta_2 \end{bmatrix}. \text{ In Matrix form, let's call } Y = A * \beta$$

By performing the linear regression, one can obtain parameters of β 's that minimizes the squared error. Therefore,

Residual_i = $y_i - A * \hat{\beta}$ where $\hat{\beta}$ is 3 by 1 matrix that minimizes the least squared error.

The results of residuals are added to the data.frame as Table.detrend. Now, the data frame is filtered based on column entries named Mn.

$$P_i = \frac{1}{\left\lfloor \frac{\text{length}(y)}{12} \right\rfloor + 1} \sum_{k=0}^{\left\lfloor \frac{\text{length}(y)}{12} \right\rfloor} y_{i+12k}$$

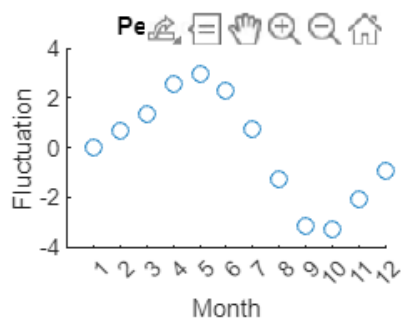
Or simply just a biased mean of samples corresponds to For y_j , $j \equiv i \pmod{12}$.

```
P_i = zeros(1,12);
for i = 1:12
    T_month = T_cutoff(T_cutoff.Mn ==i,:);
    P_i(i) = sum(T_month.detrend)/size(T_month,1);
end
scatter(1:12,P_i)
```

```

xlabel('Month')
ylabel('Fluctuation')
title('Periodic signal P_i')
xticks(1:12)

```



```

k = table(P_i');
k.Properties.VariableNames = "P";
k.Properties.RowNames =
["Jan", "Feb", "Mar", "Apr", "May", "Jun", "Jul", "Aug", "Sep", "Oct", "Nov", "Dec"];
disp(k)

```

	P
Jan	-0.012919
Feb	0.64641
Mar	1.3556
Apr	2.5619
May	2.9829
Jun	2.3165
Jul	0.7763
Aug	-1.3012
Sep	-3.1281
Oct	-3.3095
Nov	-2.0815
Dec	-0.92151

Q2: Plot the final fit

The final fit is defined as $F_n(t_i) + P_i$. Plot final fit on top of the entire time series with a vertical line indicating the split between the training and test data.

```


clear x y T_month
clf

figure()
p1 = plot(T.time,T.CO2);
x = [ones(length(T.time),1),T.time,T.time.^2];
T.extra = P_i([T.Mn])';

hold on

```

```
p2 = plot(T.time,x*b +T.extra);
xline(T.time(cutoff))
legend('Data', 'Fitted', 'training cutoff line',Location='northwest')
hold off
xlim([0 max(T.time)])
```



```
clear extra
```

Q3: Reported error of RMSE and MAPE(Max 200 words)

RMSE(Root-Mean-Square Error) is defined as following.

$$E = \sqrt{\frac{1}{n} \sum_{i=1}^n |A_i - F_i|^2}$$

MAPE(Mean Absolute Percentage Error) is defined as following.

$$E = \frac{1}{n} \sum_{i=1}^n \left| \frac{A_i - F_i}{A_i} \right| \times 100$$

Notice that error was applied only to the training set.

```
clear x
x = [ones(length(training{1}),1),training{1},training{1}.^2];
predict = x*b +T.extra(1:cutoff)
```

```
predict = 587×1
315.6238
316.8974
317.3860
315.3149
313.3054
311.5467
312.7302
313.9589
314.9363
315.6647
⋮
```

```
RMSE_Error = rmse(x*b,training{2});
```

```
MAPE_Error = mape(x*b, training{2});
fprintf('Reported RMSE and MAPE errors, (%.3f, %.3f)', RMSE_Error, MAPE_Error)
```

Reported RMSE and MAPE errors, (2.188, 0.537)

```
RMSE_Error2 = rmse(predict, training{2});
MAPE_Error2 = mape(predict, training{2});
fprintf('Reported RMSE and MAPE errors, (%.3f, %.3f)', RMSE_Error2, MAPE_Error2)
```

Reported RMSE and MAPE errors, (0.714, 0.168)

Considering seasonal variation greatly reduced errors, as they are reduced to about 1/3 of their original errors. It is intuitive, as ziggles were shown in the plot of CO₂ level; A linear fitting would ignore such ziggles.

To reason quantitatively, consider the plot of P_i where we observed seasonal variation closely resembles a sine curve.

$P_{\text{continuous}} = A \sin\left(\phi + \frac{2\pi t}{p}\right) + B$ which is a typical sine wave form introduced in class.

Then, the error generated by ignoring seasonal variation for one period would be,

$\frac{1}{N} \sum |P_{\text{continuous}}| = \frac{1}{N} \left| \sum A \sin\left(\phi + \frac{2\pi t}{p}\right) \right| \approx \frac{1}{N} \left| A \cos\left(\phi + \frac{2\pi t}{p}\right) * \frac{2\pi}{p} \right|$ sum may be approximated to integral when time intervals are sufficiently small compared to our interest. Notice by the definition of time, $p = 1$.

Also notice that having unaccounted error term produce a product of errors in RMSE:

$$(E_{\text{seasonal}} + E_{\text{Other}})^2 = (E_{\text{seasonal}})^2 + (E_{\text{Other}})^2 + 2E_{\text{seasonal}}E_{\text{Other}}$$

Therefore, having errors that are not independent might contribute to the error by forming products.

Q4: Ratio of values

1. The ratio of range of values of F to amplitude of P
2. The ratio of the amplitude of P to the range of residual R (both the trend and periodic signal)

8u7

Recall the following matrix equation of F,

$$\begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix} = \begin{bmatrix} 1 & x_1 & x_1^2 \\ 1 & x_2 & x_2^2 \\ \vdots & \vdots & \vdots \\ 1 & x_n & x_n^2 \end{bmatrix} \begin{bmatrix} \beta_0 \\ \beta_1 \\ \beta_2 \end{bmatrix}. \text{ In Matrix form, let's call } Y = A * \beta$$

In this paper, $F_n(t_i)$ is equivalent to y_i . Therefore, the range of F is $\max(Y) - \min(Y) = 74.6814$


```
P_i
```

```
P_i = 1x12
    -0.0129    0.6464    1.3556    2.5619    2.9829    2.3165    0.7763    -1.3012 ...
```

```
t = T.time(1:12)
```

```
t = 12x1
    0.2083
    0.2917
    0.3750
    0.5417
    0.6250
    0.7083
    0.8750
    0.9583
    1.0417
    1.1250
     ...
```

```
scatter(t,P_i)
```



```
beta0 = [3, 0, 1.5, 0];
modelfun = @(b,t) b(1)*sin(b(2)+2*pi.*t/b(3)) + b(4) ;
sine_fit = fitnlm(t, P_i, modelfun,beta0)
```


```
sine_fit =
Nonlinear regression model:
    y ~ b1*sin(b2 + 2*pi*t/b3) + b4
```

Estimated Coefficients:				
	Estimate	SE	tStat	pValue
b1	2.8274	0.28192	10.029	8.3061e-06
b2	-1.5152	0.2151	-7.0444	0.00010777

b3	1.1357	0.058085	19.553	4.8656e-08
b4	0.29852	0.21047	1.4183	0.19385

Number of observations: 12, Error degrees of freedom: 8
 Root Mean Squared Error: 0.679
 R-Squared: 0.928, Adjusted R-Squared 0.901
 F-statistic vs. constant model: 34.4, p-value = 6.4e-05


```
clear predict
plot(t,P_i)
hold on
x1 = linspace(t(1),t(12));
y1 = sine_fit.Coefficients.Estimate;
plot(x1, modelfun(y1,x1))
hold off
legend('P_i values', 'Non-Linear fitted')
```




```
clear x1 y1
```

From non-linear fitting, estimated parameter states that amplitude is 2.8274.

```
clear x
x = [ones(length(training{1}),1),training{1},training{1}.^2];
residual1 = training{2} - x*b - predict(sine_fit, (mod(T_cutoff.count, 12) +
0.5)/12 );
plot(residual1)
```



```
figure()
residual2 = training{2} - x*b - T.extra(1:cutoff);
plot(residual2)
```



```
fprintf('Non-linear fitting gave the range of residuals %.4f',max(residual1)-min(residual1))
```

Non-linear fitting gave the range of residuals 11.5666

```
fprintf('Subtracting the average of each months(constant) gave the range of residual %.4f',max(residual2)-min(residual2))
```

Subtracting the average of each months(constant) gave the range of residual 3.8364

It was found that the range of residuals would be smaller by subtracting P_i 's for each month. It was found that $R = 3.8364$

Rest of them are trivial arithmetics;

$$1. \frac{F}{P} = \frac{74.6814}{2.8274} \approx 26.41$$

$$2. \frac{P}{R} = \frac{2.8274}{3.8364} \approx 0.7370$$

```
F = 74.6814;
P = 2.8274;
F/P;
R = 3.8364;
P/R;
```

Decomposition of the variation of the CO_2 concentration can be meaningful based on applications. For an application to show the global trend of CO_2 for 80 years, it may not be important as $F \gg P$. However,

forecasting CO_2 concentration of a few months in the future, the seasonal variation plays a great role as $\frac{P}{\Delta F}$

and $\frac{P}{R}$ are considerable. Meaning, the change of the level from a month to another is greatly dependent on seasonal variation.

Written Exercise: ARIMA

Q1: Consider MA(1) model, find autocovariance function

Consider MA(1) model,

$y_t = W_t + \theta W_{t-1} = W_t(1 + \theta L)$ where L is a lag operator with properties $L^k X_t = X_{t-k}$

And lag polynomial is defined as following;

$$\theta(L) = 1 + \sum_{i=1}^q \theta_i L^i \quad \varphi(L) = 1 - \sum_{i=1}^p \varphi_i L^i$$

```
Mdl = arima(0,0,1)
```

```
Mdl =  
  arima with properties:  
  
    Description: "ARIMA(0,0,1) Model (Gaussian Distribution)"  
    Distribution: Name = "Gaussian"  
        P: 0  
        D: 0  
        Q: 1  
    Constant: NaN  
        AR: {}  
        SAR: {}  
        MA: {NaN} at lag [1]  
        SMA: {}  
    Seasonality: 0  
        Beta: [1x0]  
    Variance: NaN
```

```
MA_residual = estimate(Mdl,residual2)
```

ARIMA(0,0,1) Model (Gaussian Distribution):

	Value	StandardError	TStatistic	PValue
Constant	-0.00071834	0.035087	-0.020473	0.98367
MA{1}	0.69687	0.031515	22.112	2.4309e-108
Variance	0.23754	0.015796	15.037	4.1818e-51

```
MA_residual =  
  arima with properties:  
  
    Description: "ARIMA(0,0,1) Model (Gaussian Distribution)"  
    Distribution: Name = "Gaussian"  
        P: 0  
        D: 0  
        Q: 1  
    Constant: -0.000718339  
        AR: {}  
        SAR: {}  
        MA: {0.696865} at lag [1]  
        SMA: {}  
    Seasonality: 0  
        Beta: [1x0]  
    Variance: 0.237537
```

```
leftover = infer(MA_residual,residual2,'Y0',residual2(1:2));  
y_hat = residual2 - leftover
```

```
y_hat = 587x1
```

```


-0.0007
0.0529
0.3475
-0.1564
0.4882
0.7912
0.6070
-0.0057
0.4988
0.1002
⋮

```

```

plot(y_hat)
hold on
plot(residual2)
hold off
legend('Model Prediction','Residual')


```



```

figure()
[acf, ~] = autocorr(leftover, NumSTD=0);
autocorr(leftover, NumSTD=0)

```



```

[xcf, lag] = xcov(leftover);

```

This paper reports the following results;

$$y_t = W(\mu, \sigma^2)_t(1 + \theta L) = W(0, 0.2375^2)_t(1 + 0.6969 * L)$$

For autocovariance, XCF(cross correlation function) it is defined as following;

$$C_{xy}(m) = \begin{cases} \frac{1}{T} \sum_{t=1}^{T+k} (y_{1,t} - \bar{y}_1)(y_{2,t} - \bar{y}_2); & k = 0, 1, 2, \dots \\ \frac{1}{T} \sum_{t=1}^{T+k} (y_{1,t} - \bar{y}_1)(y_{2,t} - \bar{y}_2); & k = 0, -1, -2, \dots \end{cases}$$

From simple algebra, autocorr = $\frac{C_k}{C_0}$ where C was defined previously. Therefore,

$$C_k = \text{autocorr} * C_0$$

```
fprintf('Autocovariance of the first 10 lags: ')
```

Autocovariance of the first 10 lags:

```
disp(xcf(587)*acf(1:10))
```

```
139.4338
66.2124
109.7654
69.5042
91.4393
69.3541
79.0661
69.1657
72.8785
69.9372
```

Q2: AR(1) model

For AR(1) model,

$X_t = \phi X_{t-1} + W_t$ or equivalently, for the consistency of equation format,

$$(1 - \phi L)X_t = W_t$$

```
AR = arima(1,0,0)
```

AR =

arma with properties:

Description: "ARIMA(1,0,0) Model (Gaussian Distribution)"

Distribution: Name = "Gaussian"

P: 1

D: 0

Q: 0

Constant: NaN

AR: {NaN} at lag [1]

SAR: {}

MA: {}

SMA: {}

Seasonality: 0

```
Beta: [1x0]
Variance: NaN
```

```
AR_residual = estimate(AR,residual2)
```

```
ARIMA(1,0,0) Model (Gaussian Distribution):
```

	Value	StandardError	TStatistic	PValue
Constant	0.00013199	0.012331	0.010704	0.99146
AR{1}	0.90864	0.016249	55.92	0
Variance	0.088947	0.0048973	18.163	1.0215e-73


```
AR_residual =
    arima with properties:
```

```
Description: "ARIMA(1,0,0) Model (Gaussian Distribution)"
Distribution: Name = "Gaussian"
P: 1
D: 0
Q: 0
Constant: 0.00013199
AR: {0.908643} at lag [1]
SAR: {}
MA: {}
SMA: {}
Seasonality: 0
Beta: [1x0]
Variance: 0.088947
```

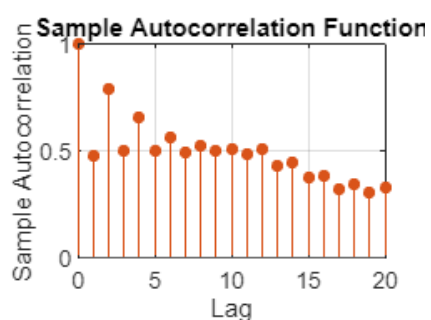
```
leftover2 = infer(AR_residual,residual2);
y_hat2 = residual2 - leftover2
```

```
y_hat2 = 587x1
    0.0632
    0.0694
    0.5022
    0.1128
    0.4954
    1.4763
    1.5114
    0.5451
    0.6463
    0.5850
    ⋮
    ⋮
```

```
plot(y_hat2)
hold on
plot(residual2)
hold off
legend('Model Prediction','Residual')
```



```
figure()
[acf2 ,~] = autocorr(leftover2,NumSTD=0);
autocorr(leftover,NumSTD=0)
```



```
[xcf2, lag2] = xcov(leftover2);
```

This paper reports the following parameters;

$$(1 - 0.908643L)X_t = W(0, 0.088947^2)_t$$

For autocovariance,

```
fprintf('Autocovariance of the first 10 lags: ')
```

Autocovariance of the first 10 lags:

```
disp(xcf2(587)*acf2(1:10))
```

```
52.2119
-9.5560
 2.0354
-1.1709
 1.9369
-0.8885
-0.3401
-0.0754
-0.5350
 3.1220
```

Written Exercise: CPI and BER

Data Explanation:

1. CPI(consumer price index)
2. BER(break-even rate)


Data loading

```
clc
clear
opts = detectImportOptions('CPI.csv');
T = readtable('CPI.csv',opts);
dim = size(T);


processed_T = rmmissing(T);
clear T
T = table2timetable(processed_T);
T.scalaritime = seconds(T.date - T.date(1))/(3600*24);
```

Q0: Verify Your Result By plotting(As recommended by Alan)

```
plot(T.date,T.CPI)
x = [ones(size(T.date,1),1), T.scalaritime];
y = T.CPI;
b = x\y;
hold on
plot(T.date,x*b)
hold off
legend('CPI','Linear Fitting')
```



```
figure()
residual = y - x*b;
plot(residual)
axis('tight')
title('residual plotted')
```

AR Model, AR model find parameters, MRSE, 1month ahead forecast

Q1:Repeat the model fitting and evaluation procedure, with rate

1: Description of calculation

Following the definition of inflation rate given by the course, we have;

$$IR_t = \frac{CPI_t - CPI_{t-1}}{CPI_{t-1}}$$

1. Find the start date and end date, which turned out to be 07/24/2008 and 10/01/2019.
2. Extract Month and year from table.date
3. Use 'Group_by' to group datas by month and year. In MATLAB, the function is called grpstats.
4. Pass Grouped datas to 'sapply' function. It is called varfun in MATLAB. Since grpstats calculates mean, it was not necessary to call.
5. Having vector(list) with CPIs in order, run a simple for loop to calculate monthly IR. And plot.

```
mth = month(T.date);
yr = year(T.date);
[monthlyCPIMean, groupName] = grpstats(T.CPI, [yr, mth], {'mean', 'gname'})
```

```
monthlyCPIMean = 136x1
```


```
100.0000
100.5251
100.1238
99.9854
98.9754
97.0797
96.0757
96.4938
96.9737
97.2095
⋮
```

```
groupName = 136x2 cell
```

```
'2008'      '7'
'2008'      '8'
'2008'      '9'
'2008'     '10'
'2008'     '11'
'2008'     '12'
'2009'      '1'
'2009'      '2'
'2009'      '3'
```


'2009'
'4'
⋮

```
plot(monthlyCPIMean)
title('IR')
```



```
IR = zeros(length(monthlyCPIMean)-1,1);
for i = 1:length(monthlyCPIMean)-1
    IR(i) = (monthlyCPIMean(i+1) - monthlyCPIMean(i))/monthlyCPIMean(i);
end

plot(IR)
title('residual of IR')
```



Q2: Describe how the data has been detrended and plot

Since regression was already performed on CPI data, I would assume that the data that the question is referring is IR data. If one is interested in detrending of CPI data, refer to Q0.

Since mathematical backgrounds of regressions were already introduced previously, without redundancy, regression would be performed. Please recall the expression given already.


$$\begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix} = \begin{bmatrix} 1 & x_1 \\ 1 & x_2 \\ \vdots & \vdots \\ 1 & x_n \end{bmatrix} \begin{bmatrix} \beta_0 \\ \beta_1 \end{bmatrix}. \text{ In Matrix form, let's call } Y = A * \beta$$

I detrended residual graph by finding inverse of matrix A after projecting Y and A into the column space of A.


AKA $(A^T A)^{-1} A^T Y = \beta$.

```
x = [ones(size(monthlyCPIMean,1),1), (1:length(monthlyCPIMean))'];
y = monthlyCPIMean;
b = x\y;
```

```
figure()
plot(monthlyCPIMean)
hold on
plot(x*b)
hold off
legend('IR', 'Linear fitting')
```



```
figure()
residual = y - x*b;
plot(residual)
axis('tight')
title('residual plotted')
```



Notice that the linear fitting seems to fail reduce the magnitude of oscillation.

Q3: Statement of and justification for the chosen AR(p) model

1. Includes: plots and reasoning

After searching for methodology online, it is recommended to use PCAF(partial Autocorrelation function) in conjunction with Yule-Walker equations. The estimation of the PACF involves solving Yule-Walker equations with respect to the autocorrelations. PCAF measures the correlation between y_t and $L^k y_t$ after adjusting for the linear effects of $L^k y_t$. L is a lag operator introduced previously.


Introducing Yule-Walker equations, the equation can be written as $\gamma_m = \sum_{k=1}^p \varphi_k \gamma_{m-k} + \sigma_e^2 \delta_{m,0}$ where $m=0, \dots, p$. To see how the equations yields $p+1$ equations, the equation was written in a matrix format.

γ_m is the autocovariance function of X_t , σ_e is the standard deviation of the input noise processes, and δ is a Kronecker delta.


$$\begin{bmatrix} \gamma_1 \\ \gamma_2 \\ \vdots \\ \gamma_p \end{bmatrix} = \begin{bmatrix} \gamma_0 & \gamma_{-1} & \gamma_{-2} & \dots \\ \gamma_1 & \gamma_0 & \gamma_{-1} & \dots \\ \gamma_2 & \gamma_1 & \gamma_0 & \dots \\ \vdots & \vdots & \vdots & \vdots \end{bmatrix} \begin{bmatrix} \varphi_1 \\ \varphi_2 \\ \vdots \\ \varphi_p \end{bmatrix} \quad \text{and} \quad \gamma_0 = \sum_{k=1}^p \varphi_k \gamma_{-k} + \sigma_e^2$$

This equation provides $p + 1$ equations for a model with $p+1$ unknowns (φ 's and σ) and none of them are redundant as each equations involves a new term. This concludes the reasoning for choosing AR(p) model.

```
parcorr(residual,Method="yule-walker", numlags = 45);
legend('PACF of residual')
```



```
parcorr(monthlyCPIMean,Method="yule-walker", numlags = 30);
legend('PACF of IR')
```



Analyzing with PCAF revealed AR(2) for residual(detrended data) and AR(1) for IR data. Therefore, one can come up with following scheme;

IR = AR(2) model + linear trend + other residual

Therefore, this paper proceeds with AR(2) model

Q4: Describe the final model;

1. Model specification
2. Plot 1month-ahead forecasts
3. Plot prediction and data together

Let's use the first 100 months for training and predict IR of 101 st month.

```
cutoff = 100
```

```
cutoff = 100
```

```
AR = arima(2,0,0);
AR_residual = estimate(AR, residual(1:cutoff))
```

ARIMA(2,0,0) Model (Gaussian Distribution):

	Value	StandardError	TStatistic	PValue
Constant	-0.017102	0.034547	-0.49503	0.62058
AR{1}	1.4948	0.059866	24.969	1.3409e-137
AR{2}	-0.62139	0.057431	-10.82	2.7739e-27
Variance	0.10064	0.013149	7.6539	1.9501e-14

AR_residual =

arma with properties:

Description: "ARIMA(2,0,0) Model (Gaussian Distribution)"

Distribution: Name = "Gaussian"

P: 2

D: 0

Q: 0

Constant: -0.0171017

AR: {1.49476 -0.621394} at lags [1 2]

SAR: {}

```

MA: {}
SMA: {}
Seasonality: 0
Beta: [1x0]
Variance: 0.10064

```

```


leftover = infer(AR_residual, residual(1:cutoff));
y_hat2 = residual(1:cutoff) - leftover;

[YF,YMSE] = forecast(AR_residual,10,residual(1:cutoff));

figure()
hold on
plot(100:109, YF, 'r')
plot(residual(1:110), 'b')
plot(y_hat2)
xline(cutoff)
plot(100:109,YF +1.96*sqrt(YMSE), 'r:')
plot(100:109,YF -1.96*sqrt(YMSE), 'r:')
k = legend('prediction', 'Residual', 'Model fit', 'cutoff', Location='northeast');
hold off

xlim([60 170])
ylim([-2.6 3.8])

```



Notice that the red dotted lines are 95 confidence intervals for forecasting.

This paper reports the results as follows;

The model found:

$$AR(2) : (1 - L\varphi_1 - L^2\varphi_2)X_t = (1 - 1.49476L + 0.621394L^2)X_t =$$

$$W(0, 0.10064^2)_t - 0.0171017$$

```
linear_predict = x*b;
```

```
fprintf('The model predicted 101th month\'s as: %f ', linear_predict(101) + YF(1))
```

The model predicted 101th month's as: 110.650578

```
fprintf('whereas the data indicates: %f ', monthlyCPIMean(101))
```

whereas the data indicates: 110.471900

Q5: Which AR(p) model to choose

1. Include a plot of the RSME against different lags p for the model

In conjunction with Yule-Walker equation, this paper stated why would PCAF give sufficient information to determine the degree of AR models.


Let's approach model selection based on RMSE this time. Like before, the training sets are 100 months of IR datas.

```
clear AR AR_residual leftover y_hat2 y_hat YF YMSE
error_mat = zeros(30,10);
% for i = 1:10
%     AR = arima(i,0,0);
%     AR_residual = estimate(AR, residual(1:cutoff));
%     YF = forecast(AR_residual,30,residual(1:cutoff));
%     error_mat(:,i) = linear_predict(100:130-1) + YF;
% end
disp(rmse(monthlyCPIMean(100:130-1),error_mat))
```

113.1257 113.1257 113.1257 113.1257 113.1257 113.1257 113.1257 113.1257 113.1257 113.1257

```
X = NaN(1,10);
X = categorical(X);
for i =1:10
    X(i) = sprintf("AR(%d)",i);
end
graph1 = bar(X,rmse(monthlyCPIMean(100:130-1),error_mat));

ylim([0.40 1.1 ])
```



From AR(1) to AR(10) models, RMSE errors are listed as followed. Notice that I predicted upto 30 months in the future, which would explain the high level of errors in general. Secondly, the RMSE error does not monotonically decreases with increasing the order since higher than necessary orders would overfit the errors in training data set and wrongly predict.

In fact, the AR(10) model has the highest RMSE error.


Q6: Overlay your estimates

1. plot IR data/predictions/BER

```
clear T2 mth yr groupName residual
T2 = readtable('T10YIE.csv');
T2 = rmmissing(T2);

mth = month(T2.DATE);
yr = year(T2.DATE);

[monthlyBERmean, groupName] = grpstats(T2.T10YIE, [yr, mth], {'mean', 'gname'});
plot(monthlyBERmean)
```



```
BER = zeros(length(monthlyBERmean)-1,1);
for i = 1:length(monthlyBERmean)-1
    BER(i) = (monthlyBERmean(i+1) - monthlyBERmean(i))/monthlyBERmean(i);
end
```



```

cutoff = 128;
x = [ones(cutoff,1), (1:cutoff)'];
y = IR(1:cutoff);
b = x\y;
residual = y - x*b

```

```

residual = 128x1
    0.0047
   -0.0045
   -0.0019
   -0.0107
   -0.0197
   -0.0109
    0.0038
    0.0044
    0.0018
    0.0019
     :
     :
     :

```

```

clear linear_predict
linear_predict = x*b

```

```

linear_predict = 128x1
    0.0005
    0.0005
    0.0005
    0.0006
    0.0006
    0.0006
    0.0006
    0.0006
    0.0006
    0.0006
    0.0006
     :
     :
     :

```

```

AR = arima(2,0,0);
AR_residual = estimate(AR, residual)

```

ARIMA(2,0,0) Model (Gaussian Distribution):

	Value	StandardError	TStatistic	PValue
Constant	3.0882e-05	0.00028805	0.10721	0.91462
AR{1}	0.65816	0.071846	9.1608	5.1528e-20
AR{2}	-0.24811	0.070262	-3.5312	0.00041364
Variance	9.33e-06	5.1659e-07	18.061	6.4957e-73

AR_residual =

arima with properties:

Description: "ARIMA(2,0,0) Model (Gaussian Distribution)"

Distribution: Name = "Gaussian"

P: 2

D: 0

Q: 0

Constant: 3.08817e-05

AR: {0.658164 -0.24811} at lags [1 2]

```

SAR: {}
MA: {}
SMA: {}
Seasonality: 0
Beta: [1x0]
Variance: 9.32997e-06

```

```

leftover = infer(AR_residual, residual);
y_hat2 = residual - leftover;
[YF,~] = forecast(AR_residual,135-cutoff,residual);

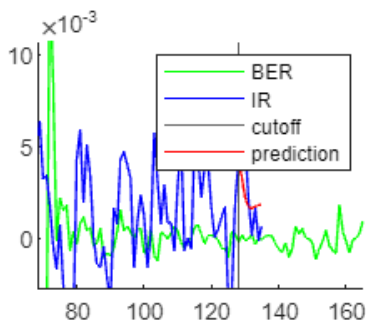
```

```

figure()
% plot(BER)
hold on
% plot(BER)
plot(BER/100, 'g')
plot(IR, 'b')
xline(cutoff)
clear x
x = [ones(7,1), (cutoff+1:length(IR))'];
plot(cutoff:length(IR), [IR(cutoff);(x*b+YF)], 'r')
hold off
legend('BER', 'IR', 'cutoff', 'prediction')

xlim([68.7 165.0])
ylim([-0.0027 0.0108])

```




As I do not understand what three graphs were supposed to be plotted, I plotted BER rate, IR rate, and AR(2) model fitted by IR.

```

stem(xcov(IR,BER))
xlim([124 244])
ylim([-0.0074 0.0118])

```




Since I didn't see a clear pattern between IR rate and BER rate, I plotted covariance between them. It seems like there are a seasonal pattern hiding.

Written Assignment #3

Q1: Plot the cross correlation function between the and inflation rate

```
xdata = xcorr(IR,BER);
stem(xdata)

xlim([186.6 254.1])
ylim([-0.0031 0.0077])
```



What is shown is a magnified portion of cross-correlation function between IR and BER. For me, the lag term is not super concise, but I occasionally counted 12 terms before sin-like curve ends. And it does make sense to have a strong seasonal lag with 12, such as black friday sale and so on.

Q2: Fit SARIMAX

```
data = [IR(1:cutoff),BER(1:cutoff)];

clear AR AR_residual

AR = SARIMAX_data1;
```

Unrecognized function or variable 'SARIMAX_data1'.

```
AR_residual = estimate(AR, IR)
leftover = infer(AR_residual, IR);
y_hat2 = IR - leftover;
```

```
[YF,~] = forecast(AR_residual,5,IR);
```

```
cutoff = 128;
x = [ones(cutoff,1), (1:cutoff)'];
y = IR(1:cutoff);
b = x\y;
residual = y - x*b
clear linear_predict
linear_predict = x*b

figure()
hold on
plot(IR, 'b')
xline(cutoff)
clear x
x = [ones(5,1), (cutoff:cutoff+4)'];
plot(cutoff:cutoff+5, [IR(cutoff);(x*b+YF)], 'r')
hold off
legend('IR', 'cutoff', 'prediction')

xlim([88.1 175.1])
ylim([-0.0063 0.0123])
```

This papaer reports SARIMAX model of;

$$(1 - \phi_1 L - \phi_2 L^2)(1 - L^{12})y_t = c + X_1 \beta_1 + \epsilon_t =$$

$$(1 - 0.672424L + 0.289171L^2)(1 - L^{12})y_t = 0.00082287 - 0.000870617X_1 + \epsilon(0, (8.4e - 06)^2)_t$$

For prediction;

```
fprintf('the prediction is %f', (linear_predict(1)+YF(1)))
fprintf('whereas the actual data is %f', IR(101))
```

Q3: RMSE

This paper reports rmse between IR data and SARIMAX model to be;

```
rmse_data = vertcat(y_hat2(1:cutoff), x*b+YF)
fprintf('The RMSE error is : %.10f', rmse(IR(1:133),rmse_data))
```

Q4: Improving Your Model


1. What are other steps to improve?
2. What is the smallest error you can obtain?

3. Describe the model that performs best

There were several arima parameters that I have not considered. The arima models allows lags for X_t , ϵ_t , and degree of integration, along with seasonal parameters with exogenous parameter called betas.

As I am not fully comprehend the dynamics of parameters, I wasn't capable of exploring. But I would imagine that with sufficient trial and errors, one can come up with a better model if such parameters are adequately investigated.

After trying several models, following model came out to be the most successful.



```
clear AR AR_residual
AR = SARIMAX_data1;
AR_residual = estimate(AR, IR)
leftover = infer(AR_residual, IR);
y_hat2 = IR - leftover;
fprintf('RMSE error: %.10f', rmse(y_hat2,IR))
```